

```
In [ ]: import requests
        from PIL import Image
        import io
        # import tensorflow as tf
        import torch
        import datetime
        from pathlib import Path
        from datasets import DatasetDict
        from datasets import Dataset
        from datasets import load_from_disk
        import time
        from tqdm import tqdm
```

```
In [ ]: from datasets import load_dataset

        # this line is just for quick testin
        # dataset = load_dataset("stochastic/random_streetview_images_pano_v0.0.2", split="train[:5000]")
        dataset = load_dataset("stochastic/random_streetview_images_pano_v0.0.2", split="train")
```

```
In [ ]: # filter all the countries with not many images out of the dataset
        passes = ['HK', 'HU', 'AU', 'SK', 'FR', 'FI', 'ZA', 'LT', 'PE', 'US'] # 10 most common

        # Function to apply the filter
        def filter_countries(example):
            return example['country_iso_alpha2'] in passes

        # Apply the filter to the dataset
        dataset = dataset.filter(filter_countries)
```

```
Filter: 100%|██████████| 11054/11054 [07:44<00:00, 23.81 examples/s]
```

```
In [ ]: # test out the filtering
        dataset
```

```
Out[ ]: Dataset({
  features: ['image', 'country_iso_alpha2', 'latitude', 'longitude', 'address'],
  num_rows: 2151
})
```

```
In [ ]: # preprocessing - this took 6min 48 seconds on my computer... jk 12 min second time
        from torchvision.transforms import v2
        import torchvision.transforms as transforms
```

```

# data preprocessing
image_transforms = v2.Compose([
    v2.CenterCrop((561, 1010)), # crops the middle image from the panorama
    v2.Resize((224, 224)), # resizing the image to 224x224 for easier processing
    transforms.ToTensor() # converting the image to a tensor
])

def transform(batch):
    # Transform each image in the batch and ensure it has 3 channels
    batch["image"] = [image_transforms(image.convert("RGB")) for image in batch["image"]]
    del batch["latitude"]
    del batch["longitude"]
    del batch["address"]
    return batch

# Apply the transformations to the dataset
dataset = dataset.map(transform, batched=True, batch_size=8)

```

```
Map: 100%|██████████| 2151/2151 [01:34<00:00, 22.75 examples/s]
```

```

In [ ]: # check what the name of the country thing is
        # print(dataset[0])

```

```

In [ ]: # Make split and save
train_testvalid = dataset.train_test_split(test_size=0.4)
test_valid = train_testvalid['test'].train_test_split(test_size=0.5)

datasets = DatasetDict({
    'train': train_testvalid['train'],
    'test': test_valid['test'],
    'valid': test_valid['train']
})
datasets.save_to_disk("./data")

```

```

Saving the dataset (0/2 shards): 0%|██████████| 0/1290 [00:00<?, ? examples/s]
Saving the dataset (2/2 shards): 100%|██████████| 1290/1290 [00:03<00:00, 405.90 examples/s]
Saving the dataset (1/1 shards): 100%|██████████| 431/431 [00:04<00:00, 98.55 examples/s]
Saving the dataset (1/1 shards): 100%|██████████| 430/430 [00:04<00:00, 86.71 examples/s]

```

```

In [ ]: # country_codes = ["ZA", "KR", "AR", "BW", "GR", "SK", "HK", "NL", "PE", "AU", "KH", "LT", "NZ", "RO", "MY", "SG", "AE", "FR", "ES", "IT",
        country_codes = ['HK', 'HU', 'AU', 'SK', 'FR', 'FI', 'ZA', 'LT', 'PE', 'US']
        country_dict = {}
        # TODO: these might need to be tensor arrays but thats easy enough to change if needed
        for i in range(len(country_codes)):

```

```

country_dict[country_codes[i]] = [0.]*len(country_codes)
country_dict[country_codes[i]][i] = 1.
# print(country_dict)
country_dict_not_one_hot = {}
for i in range(len(country_codes)):
    country_dict_not_one_hot[country_codes[i]] = i

```

```

In [ ]: # referenced: https://blog.paperspace.com/convolutional-autoencoder/
# autoencoder classes (CREDIT: LARGELY TAKEN FROM 6_AUTOENCODER NOTEBOOK, but encoder and decoder architectures modified)
# should only have Encoder that has a latent dimension of 50 - corresponding to country weights
import torch.nn as nn
import torch.nn.functional as F

class MLPEncoder(torch.nn.Module):

    def __init__(self,
                  number_of_hidden_layers: int,
                  latent_size: int,
                  hidden_size: int,
                  input_size: int,
                  activation: torch.nn.Module):

        super().__init__()

        self.latent_size = latent_size
        assert number_of_hidden_layers >= 0, "Decoder number_of_hidden_layers must be at least 0"

        # Convolutional layers
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, stride=2, padding=1)
        self.conv2 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=2, padding=1)
        self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=2, padding=1)
        self.conv4 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, stride=2, padding=1)

        self.feature_map_size = 128 * 14 * 14

        # Fully connected layer to produce the latent representation of size 50
        self.fc = nn.Linear(self.feature_map_size, latent_size)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = x.view(-1, self.feature_map_size) # Flatten the output

```

```
x = self.fc(x)
return x
```

```
In [ ]: # define our training parameters and model
hidden_layers = 4
hidden_size = 30
latent_size = 10
input_size = 224
lr = 0.005
lamb = 1

# fix random seed
torch.manual_seed(0)

# select device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = MLPDecoder( number_of_hidden_layers=hidden_layers,
                    latent_size=latent_size,
                    hidden_size=hidden_size,
                    input_size=input_size,
                    activation=torch.nn.ReLU()).to(device)

# use an optimizer to handle parameter updates
opt = torch.optim.Adam(model.parameters(), lr=lr, weight_decay=1e-5)

# save all log data to a local directory
run_dir = "logs"

# to clear out TensorBoard and start totally fresh, we'll need to
# remove old logs by deleting them from the directory
!rm -rf ./logs/

# timestamp the logs for each run so we can sort through them
run_time = datetime.datetime.now().strftime("%I%M%p on %B %d, %Y")

# initialize a SummaryWriter object to handle all logging actions
from torch.utils.tensorboard import SummaryWriter
logger = SummaryWriter(log_dir=Path(run_dir) / run_time, flush_secs=20)
```

```
In [ ]: # model(data['train'][0]['image']).to(device)).shape
```

```
In [ ]: # load data from disk
data = load_from_disk("./data")
data = data.with_format("torch")
```

```
In [ ]: # test the datatype of the dataset – the datasets['train']['image'] should be a tensor
# type(datasets['train'][0]['image'])
# datasets['train'][0]['image']

torch.tensor(data['train'][0]['image']).shape
# print(len(datasets['train'][0]['image'][0]))
# print(len(datasets['train'][0]['image'][0][0]))
```

/var/folders/vc/7kh7f0_n749bhl6c8b17t3q40000gn/T/ipykernel_16242/3299156962.py:5: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

```
torch.tensor(data['train'][0]['image']).shape
```

```
Out [ ]: torch.Size([3, 224, 224])
```

```
In [ ]: # (**credit**): mostly taken from provided notebook )
# training
torch.autograd.set_detect_anomaly(False)

epochs = 25
batch_size = 1000
start_time = time.time()
loss_history = []
valid_history = []
acc_history = []
valid_acc_history = []
report_every = 1
# Loss = torch.nn.BCELoss()
Loss = torch.nn.CrossEntropyLoss()

train_loader = torch.utils.data.DataLoader(data['test'], batch_size=batch_size, shuffle=True)
valid_loader = torch.utils.data.DataLoader(data['valid'], batch_size=batch_size, shuffle=True)

for epoch in range(epochs):
    model.train()
    # weight batch losses/scores proportional to batch size
    iter_count = 0
    valid_iter_count = 0
    loss_epoch = 0.
    class_accuracy_epoch = 0.
    valid_loss_epoch = 0.
    valid_accuracy_epoch = 0.
    ###
    ### IMAGE_DATA_TRAIN is training data, shape is 610 x 3 x 64 x 64
```

```
###
```

```
# print(batched_image_data_train[0][0].shape)
#batched_image_data_train = [batched_image_data_train[0][0].unsqueeze(0)]
for idx, batch in enumerate(train_loader):

    x = batch['image']
    x = torch.tensor(x)
    # print(x)
    # flatten input images and move to device
    # *****
    # x = x / 255
    # plot x_real later to see if this is correct
    x = x.to(device)
    model.zero_grad()
    opt.zero_grad()

    # train on a batch of inputs
    pred_labels = model(x)
    # print(pred_labels)

    # get the true label
    # label = torch.tensor(country_dict[img['country_iso_alpha2']], dtype=torch.float).to(device).unsqueeze(0)
    # print(batch['country_iso_alpha2'])
    labels = torch.tensor([country_dict_not_one_hot[country] for country in batch['country_iso_alpha2']]).to(device)
    # print(labels)
    # print(len(country_dict))
    # pred labels is 4x56
    # print(len(pred_labels))
    # print(len(pred_labels[0]))
    # label is 1x56

    loss = Loss(pred_labels, labels)
    loss.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    opt.step()

    # log loss
    loss_epoch += loss.detach().item()

    # classification accuracy
    # add 1 to class_accuracy_epoch if the classification is correct, else 0
    # find index of max probability from pred_labels
    for i in range(len(pred_labels)):
        c = torch.argmax(pred_labels[i])
```

```

        # find index of 1 from label
        true_class = torch.argmax(labels[i])
        acc = (true_class == c)
        class_accuracy_epoch += acc
        iter_count += 1
    c = torch.argmax(pred_labels)

    # print(f"true: {true_class}")
    # print(f"pred: {classification.mean()}")
    # print(f"acc: {acc}")

    # plot loss
    class_accuracy_epoch /= iter_count #accuracy as a percent
    # print(iter_count)
    # logger.add_scalar("mse_loss", loss_epoch, epoch)
    loss_history.append(loss_epoch)
    acc_history.append(class_accuracy_epoch)

    # logger.add_scalar("mse_loss_valid", valid_loss_epoch, epoch)
    # # plot example generated images
    # with torch.no_grad():
    #     reconstructed_batch = model(example_batch.reshape(batch_size, -1)).reshape(batch_size, 1, image_size, image_size)
    #     logger.add_image("reconstructed_images", make_grid(reconstructed_batch, math.floor(math.sqrt(batch_size))))
    # calculate validation loss

with torch.no_grad():
    model.eval()
    for idx, valid_data in enumerate(valid_loader):
        x_valid = torch.tensor(valid_data['image']) #.float()
        # x_valid = x_valid / 255
        x_valid = x_valid.to(device)
        # print(x_valid)
        pred_labels = model(x_valid)

        labels_valid = torch.tensor([country_dict_not_one_hot[country] for country in valid_data['country_iso_alpha_2']])
        # print(label_valid)
        # print(pred_labels)
        valid_loss = Loss(pred_labels, labels_valid)
        valid_loss_epoch += valid_loss.detach().item()

    # classification accuracy
    # add 1 to class_accuracy_epoch if the classification is correct, else 0
    for i in range(len(pred_labels)):

```

```

        c = torch.argmax(pred_labels[i])
        # find index of 1 from label
        true_class = torch.argmax(labels_valid[i])
        acc = (true_class == c)
        valid_accuracy_epoch += acc
        valid_iter_count += 1
        # print(f"true: {true_class}")
        # print(f"pred: {classification.mean()}")
        # print(f"acc: {valid_acc}")
    valid_loss_epoch
    valid_history.append(valid_loss_epoch)
    valid_accuracy_epoch /= valid_iter_count
    valid_acc_history.append(valid_accuracy_epoch)

if (epoch + 1) % report_every == 0:
    mins = (time.time() - start_time) / 60
    print(f"Epoch: {epoch + 1:5d}\tMSE Loss: {loss_epoch :6.4f}\t in {mins:5.1f}min")
    print()

```

/var/folders/vc/7kh7f0_n749bhl6c8b17t3q40000gn/T/ipykernel_16242/227957138.py:38: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

```
x = torch.tensor(x)
```

/var/folders/vc/7kh7f0_n749bhl6c8b17t3q40000gn/T/ipykernel_16242/227957138.py:105: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).


```
x_valid = torch.tensor(valid_data['image'])#.float()
```


Epoch:	1	MSE Loss: 2.3009	in	0.6min
Epoch:	2	MSE Loss: 5.9839	in	1.3min
Epoch:	3	MSE Loss: 2.6905	in	1.7min
Epoch:	4	MSE Loss: 2.2884	in	2.2min
Epoch:	5	MSE Loss: 2.2844	in	2.6min
Epoch:	6	MSE Loss: 2.2758	in	3.0min
Epoch:	7	MSE Loss: 2.2711	in	3.5min
Epoch:	8	MSE Loss: 2.2611	in	4.0min
Epoch:	9	MSE Loss: 2.2473	in	4.5min
Epoch:	10	MSE Loss: 2.2168	in	4.9min
Epoch:	11	MSE Loss: 2.1905	in	5.3min
Epoch:	12	MSE Loss: 2.1737	in	5.7min
Epoch:	13	MSE Loss: 2.1363	in	6.2min
Epoch:	14	MSE Loss: 2.2527	in	6.6min
Epoch:	15	MSE Loss: 2.2335	in	7.0min
Epoch:	16	MSE Loss: 2.1195	in	7.5min
Epoch:	17	MSE Loss: 2.1297	in	7.9min
Epoch:	18	MSE Loss: 2.1363	in	8.3min
Epoch:	19	MSE Loss: 2.0965	in	8.8min
Epoch:	20	MSE Loss: 2.0707	in	9.2min
Epoch:	21	MSE Loss: 2.0451	in	9.6min
Epoch:	22	MSE Loss: 2.0142	in	10.0min
Epoch:	23	MSE Loss: 1.9965	in	10.4min


Epoch: 24 MSE Loss: 1.9765 in 10.8min

Epoch: 25 MSE Loss: 1.9325 in 11.3min

```
In [ ]: # plot loss history
import matplotlib.pyplot as plt
plt.plot(loss_history, label="training loss")
plt.plot(valid_history, label="validation loss")
plt.legend()
plt.show()
```



```
In [ ]: # plot accuracy history
plt.plot(acc_history, label="training accuracy")
plt.plot(valid_acc_history, label="validation accuracy")
plt.legend()
plt.show()
```



```
In [ ]: def predict(i):
        prediction = model(data['train'][i]['image'].to(device))
        country = torch.argmax(prediction)
        return country_codes[country]
```

```
In [ ]: from collections import Counter
        Counter([predict(i) for i in range(100)])
```

```
Out[ ]: Counter({'LT': 28,
                'HU': 17,
                'HK': 17,
                'US': 11,
                'AU': 9,
                'FR': 9,
                'ZA': 4,
                'SK': 3,
                'FI': 2})
```

```
In [ ]: # find most common countries in dataset
        from collections import Counter
```

```
country_counts = Counter(data['test']['country_iso_alpha2'])  
print(country_counts.most_common(10))
```

```
[('HU', 54), ('HK', 51), ('LT', 51), ('AU', 47), ('FR', 44), ('US', 42), ('SK', 41), ('ZA', 39), ('PE', 39), ('FI', 23)]
```

In []: